

Large-scale Logistic Regression and Linear Support Vector Machines Using Spark

Ching-Pei Lee

National Taiwan University

University of Illinois



Joint work with

Chieh-Yen Lin, Cheng-Hao Tsai and Chih-Jen Lin

IEEE 2014 Conference on Big Data, October 28, 2014

Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design
- 4 Related Works
- 5 Discussions and Conclusions



Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design
- 4 Related Works
- 5 Discussions and Conclusions



Linear Classification on One Computer

- Linear classification on one machine is a mature technique: millions of data can be trained in a few seconds.



Linear Classification on One Computer

- Linear classification on one machine is a mature technique: millions of data can be trained in **a few seconds**.
- What if the data are even bigger than the capacity of our machine?



Linear Classification on One Computer

- Linear classification on one machine is a mature technique: millions of data can be trained in **a few seconds**.
- What if the data are even bigger than the capacity of our machine?
- Solution 1: get a machine with larger memory/disk.



Linear Classification on One Computer

- Linear classification on one machine is a mature technique: millions of data can be trained in **a few seconds**.
- What if the data are even bigger than the capacity of our machine?
- Solution 1: get a machine with larger memory/disk.
 - The **data loading time** would be too lengthy.



Linear Classification on One Computer

- Linear classification on one machine is a mature technique: millions of data can be trained in **a few seconds**.
- What if the data are even bigger than the capacity of our machine?
- Solution 1: get a machine with larger memory/disk.
 - The **data loading time** would be too lengthy.
- Solution 2: **distributed training**.



Distributed Linear Classification

- In distributed training, **data loaded in parallel** to reduce the I/O time.
- With more machines, computation is faster.



Distributed Linear Classification

- In distributed training, **data loaded in parallel** to reduce the I/O time.
- With more machines, computation is faster.
- But **communication** and **synchronization** cost become significant.
- To keep the training efficiency, we need to consider **algorithms** with less communication cost, and examine **implementation details** carefully.



Distributed Linear Classification on Apache Spark

- We train logistic regression (**LR**) and L2-loss linear support vector machine (**SVM**) models on **Apache Spark** (Zaharia et al., 2010).



Distributed Linear Classification on Apache Spark

- We train logistic regression (**LR**) and L2-loss linear support vector machine (**SVM**) models on **Apache Spark** (Zaharia et al., 2010).
- Why Spark?



Distributed Linear Classification on Apache Spark

- We train logistic regression (**LR**) and L2-loss linear support vector machine (**SVM**) models on **Apache Spark** (Zaharia et al., 2010).
- Why Spark?
 - MPI (Snir and Otto, 1998) is efficient, but does not support **fault tolerance**.



Distributed Linear Classification on Apache Spark

- We train logistic regression (**LR**) and L2-loss linear support vector machine (**SVM**) models on **Apache Spark** (Zaharia et al., 2010).
- Why Spark?
 - MPI (Snir and Otto, 1998) is efficient, but does not support **fault tolerance**.
 - MapReduce (Dean and Ghemawat, 2008) supports fault tolerance, but is slow in communication.



Distributed Linear Classification on Apache Spark (cont'd)

- Why Spark?
 - Spark combines advantages of both frameworks.



Distributed Linear Classification on Apache Spark (cont'd)

- Why Spark?
 - Spark combines advantages of both frameworks.
 - Communications conducted in-memory.
 - Supports fault tolerance.



Distributed Linear Classification on Apache Spark (cont'd)

- Why Spark?
 - Spark combines advantages of both frameworks.
 - Communications conducted in-memory.
 - Supports fault tolerance.
- However, Spark is new and still **under development**.
- We therefore need to examine important **implementation issues** to ensure efficiency.



Apache Spark

- Only the **master-slave** framework.



Apache Spark

- Only the **master-slave** framework.
- **Data fault tolerance**: Hadoop Distributed File System (Borthakur, 2008).



Apache Spark

- Only the **master-slave** framework.
- **Data fault tolerance**: Hadoop Distributed File System (Borthakur, 2008).
- **Computation fault tolerance**:



Apache Spark

- Only the **master-slave** framework.
- **Data fault tolerance**: Hadoop Distributed File System (Borthakur, 2008).
- **Computation fault tolerance**: **Read-only** Resilient Distributed Datasets (RDD) and lineage (Zaharia et al., 2012).

Basic idea: reconduct operations recorded in lineage on immutable RDDs.



Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design
- 4 Related Works
- 5 Discussions and Conclusions



Logistic Regression and Linear Support Vector Machine

- Given training instances $\{(y_i, \mathbf{x}_i)\}_{i=1}^I$, $y_i \in \{-1, 1\}$, $\mathbf{x}_i \in \mathbf{R}^n$.
- Linear classification: given $C > 0$,

$$\min_{\mathbf{w}} f(\mathbf{w}) \equiv \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^I \xi(\mathbf{w}; \mathbf{x}_i, y_i)$$

$$\xi_{\text{SVM}}(\mathbf{w}; \mathbf{x}_i, y_i) \equiv \max(0, 1 - y_i \mathbf{w}^T \mathbf{x}_i)^2 \quad \text{and}$$

$$\xi_{\text{LR}}(\mathbf{w}; \mathbf{x}_i, y_i) \equiv \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i})$$

- We use a **trust region Newton method** to minimize $f(\mathbf{w})$ (Lin and Moré, 1999).



Trust Region Newton Method

- At iteration t , given iterate \mathbf{w}^t and trust region $\Delta_t > 0$, solve

$$\min_{\|\mathbf{d}\| \leq \Delta_t} q_t(\mathbf{d}) \equiv \nabla f(\mathbf{w}^t)^T \mathbf{d} + \frac{1}{2} \mathbf{d}^T \nabla^2 f(\mathbf{w}^t) \mathbf{d}$$

- $\rho_t = \frac{f(\mathbf{w}^t + \mathbf{d}) - f(\mathbf{w}^t)}{q_t(\mathbf{d})}$.
- $\mathbf{w}^{t+1} = \begin{cases} \mathbf{w}^t + \mathbf{d} & \text{if } \rho_t > \eta, \\ \mathbf{w}^t & \text{if } \rho_t \leq \eta. \end{cases}$
- Adjust the trust region size by ρ_t .
- If n is large: $\nabla^2 f(\mathbf{w}^t) \in \mathbf{R}^{n \times n}$ is too large to store.
- Consider Hessian-free methods.



Trust Region Newton Method (cont'd)

- Use a **conjugate gradient (CG) method**.
- CG is an iterative method: only needs $\nabla^2 f(\mathbf{w}^t)\mathbf{v}$ for some $\mathbf{v} \in \mathbf{R}^n$ at each iteration.
- For LR and SVM, at each CG iteration we compute

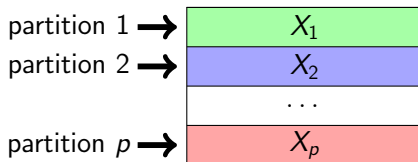
$$\nabla^2 f(\mathbf{w}^t)\mathbf{v} = \mathbf{v} + C (X^T (D (X\mathbf{v}))) , \text{ where } X \equiv \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_l \end{bmatrix}$$

is the data matrix and D is a diagonal matrix with values determined by \mathbf{w}^t .



Distributed Hessian-vector Products

- Data matrix X is **distributedly stored**

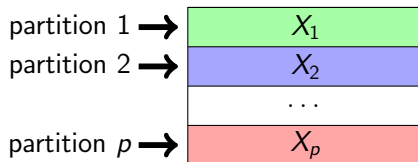


$$X^T D X \mathbf{v} = X_1^T D_1 X_1 \mathbf{v} + \dots + X_p^T D_p X_p \mathbf{v}$$



Distributed Hessian-vector Products

- Data matrix X is **distributedly stored**



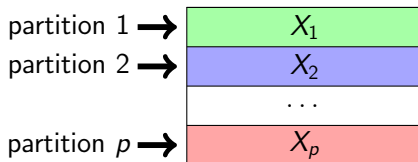
$$X^T DX\mathbf{v} = X_1^T D_1 X_1 \mathbf{v} + \dots + X_p^T D_p X_p \mathbf{v}$$

- $p \geq (\text{\#slave nodes})$ for parallelization.
- Two communications** per operation:
 - Master sends \mathbf{w}^t and the current \mathbf{v} to the slaves.
 - Slaves return $X_i^T D_i X_i \mathbf{v}$ to master.



Distributed Hessian-vector Products

- Data matrix X is **distributedly stored**



$$X^T DX\mathbf{v} = X_1^T D_1 X_1 \mathbf{v} + \dots + X_p^T D_p X_p \mathbf{v}$$

- $p \geq (\text{\#slave nodes})$ for parallelization.
- Two communications** per operation:
 - Master sends \mathbf{w}^t and the current \mathbf{v} to the slaves.
 - Slaves return $X_i^T D_i X_i \mathbf{v}$ to master.
- The same scheme for computing function/gradient.



Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design**
- 4 Related Works
- 5 Discussions and Conclusions



Experimental Settings

- We evaluate the performance by the relative difference to the optimal function value:

$$\left| \frac{f(\mathbf{w}) - f(\mathbf{w}^*)}{f(\mathbf{w}^*)} \right|.$$

- All the experiments use $C = 1$.
- We present LR results here.



Data Information

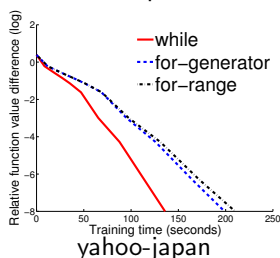
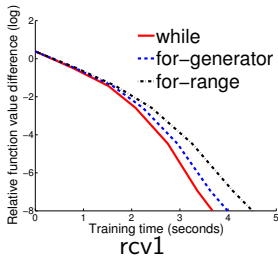
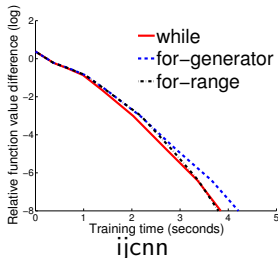
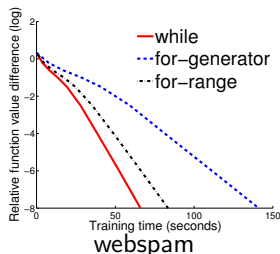
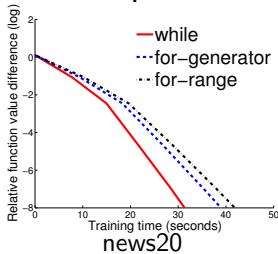
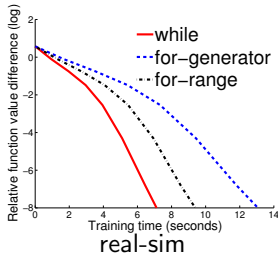
Density: avg. ratio of non-zero features per instance.

Data set	#instances	#features	density	#non-zeros
real-sim	72,309	20,958	0.25%	3,709,083
news20	19,996	1,355,191	0.03%	9,097,916
webspam	350,000	254	33.52%	29,796,333
ijcnn	49,990	22	59.09%	649,870
rcv1	20,242	47,236	0.16%	1,498,952
yahoo-japan	176,203	832,026	0.02%	23,506,415
yahoo-korea	460,554	3,052,939	0.01%	156,436,656
covtype	581,012	54	22.00%	6,901,775
epsilon	400,000	2,000	100.00%	800,000,000
rcv1t	677,399	47,236	0.16%	49,556,258



Scala Issue: Loop structures

We use one node in this experiment.



RDD: **map** or **mapPartitions**

- The second term of the Hessian-vector product

$$\sum_{i=1}^l \mathbf{x}_i D_{i,i} \mathbf{x}_i^T \mathbf{v} = \sum_{i=1}^l a(y_i, \mathbf{x}_i, \mathbf{w}, \mathbf{v}) \mathbf{x}_i,$$

where $a(y_i, \mathbf{x}_i, \mathbf{w}, \mathbf{v}) = D_{i,i} \mathbf{x}_i^T \mathbf{v}$, can be computed by either **map** or **mapPartitions**.



RDD: **map** or **mapPartitions**

- The second term of the Hessian-vector product

$$\sum_{i=1}^l \mathbf{x}_i D_{i,i} \mathbf{x}_i^T \mathbf{v} = \sum_{i=1}^l a(y_i, \mathbf{x}_i, \mathbf{w}, \mathbf{v}) \mathbf{x}_i,$$

where $a(y_i, \mathbf{x}_i, \mathbf{w}, \mathbf{v}) = D_{i,i} \mathbf{x}_i^T \mathbf{v}$, can be computed by either **map** or **mapPartitions**.

Algorithm 2 **map** implementation

```
1: data.map(new Function() {  
2:   call(y, x) { return a(y, x, w, v)x }  
3: }).reduce(new Function() {  
4:   call(a, b) { return a + b }  
5: })
```



RDD : map or mapPartitions (cont'd)

Algorithm 3 mapPartitions implementation

```
1: data.mapPartitions(new Function() {  
2:   call(partition) {  
3:     partitionHv = new DenseVector(n)  
4:     for each (y, x) in partition  
5:       partitionHv +=  $a(y, \mathbf{x}, \mathbf{w}, \mathbf{v})\mathbf{x}$   
6:   }  
7: }).reduce(new Function() {  
8:   call(a, b) { return  $\mathbf{a} + \mathbf{b}$  }  
9: })
```



RDD : map or mapPartitions (cont'd)

Algorithm 4 mapPartitions implementation

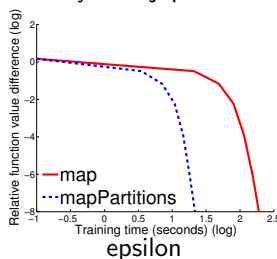
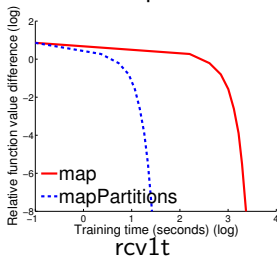
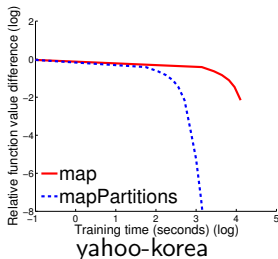
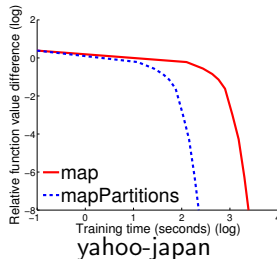
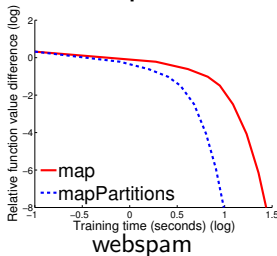
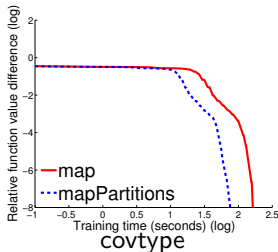
```
1: data.mapPartitions(new Function() {  
2:   call(partition) {  
3:     partitionHv = new DenseVector(n)  
4:     for each (y, x) in partition  
5:       partitionHv +=  $a(y, \mathbf{x}, \mathbf{w}, \mathbf{v})\mathbf{x}$   
6:   }  
7: }).reduce(new Function() {  
8:   call(a, b) { return  $\mathbf{a} + \mathbf{b}$  }  
9: })
```

- **map**: / **sparse** intermediate vectors.
- **mapPartitions**: **p dense** intermediate vectors.



RDD: map or mapPartitions (cont'd)

We use 16 nodes in this experiment.



Communication

- Master to slaves: Spark by default send \mathbf{w}^t and \mathbf{v} to each **partition**.
- The same \mathbf{w}^t is sent **repeatedly** in CG.



Communication

- Master to slaves: Spark by default send \mathbf{w}^t and \mathbf{v} to each **partition**.
- The same \mathbf{w}^t is sent **repeatedly** in CG.
- Use **broadcast variables** to improve.
 - Read-only variables **shared among partitions** in the same node.
 - **Cached** in the slave machines.



Communication

- Master to slaves: Spark by default send \mathbf{w}^t and \mathbf{v} to each **partition**.
- The same \mathbf{w}^t is sent **repeatedly** in CG.
- Use **broadcast variables** to improve.
 - Read-only variables **shared among partitions** in the same node.
 - **Cached** in the slave machines.
- Slaves to master: Spark by default collect results from each partition **separately**.



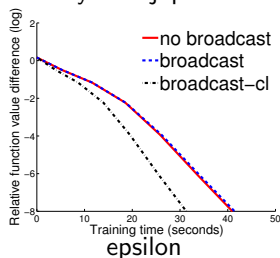
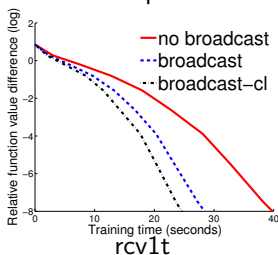
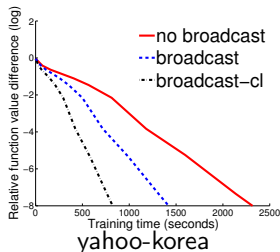
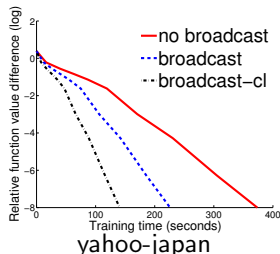
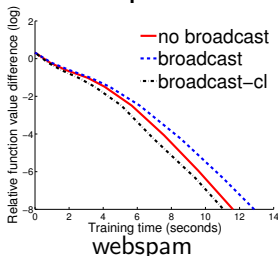
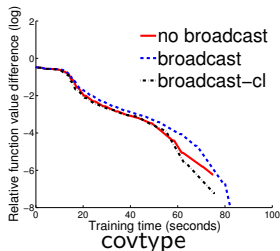
Communication

- Master to slaves: Spark by default send \mathbf{w}^t and \mathbf{v} to each **partition**.
- The same \mathbf{w}^t is sent **repeatedly** in CG.
- Use **broadcast variables** to improve.
 - Read-only variables **shared among partitions** in the same node.
 - **Cached** in the slave machines.
- Slaves to master: Spark by default collect results from each partition **separately**.
- Use the **coalesce** function.
 - **Merge** partitions on the same node before communication.



Broadcast Variables and coalesce

We use 16 nodes in this experiment.



Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design
- 4 Related Works**
- 5 Discussions and Conclusions



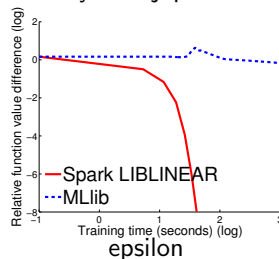
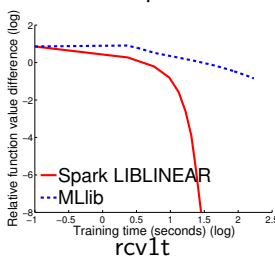
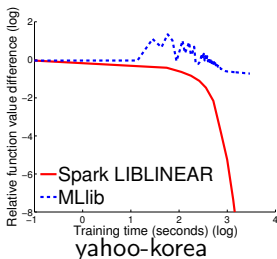
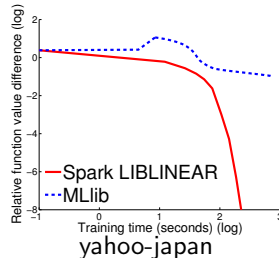
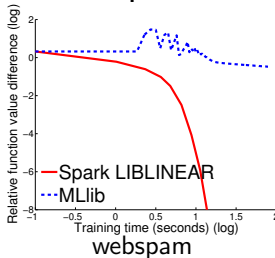
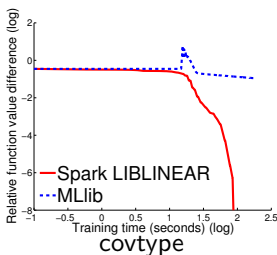
MLlib in Spark

- MLlib is a machine learning library implemented in Apache Spark.
- A stochastic gradient method for LR and SVM (but default batch size is the whole data).



Comparison with MLlib

We use 16 nodes in this experiment.



MPI LIBLINEAR

- A C++/MPI implementation by Zhuang et al. (2014) of the distributed trust region Newton algorithm we discussed.



MPI LIBLINEAR

- A C++/MPI implementation by Zhuang et al. (2014) of the distributed trust region Newton algorithm we discussed.
- No fault tolerance.
- Should be faster than our implementation:



MPI LIBLINEAR

- A C++/MPI implementation by Zhuang et al. (2014) of the distributed trust region Newton algorithm we discussed.
- No fault tolerance.
- Should be faster than our implementation:
 - More computational efficient: implemented in C++.



MPI LIBLINEAR

- A C++/MPI implementation by Zhuang et al. (2014) of the distributed trust region Newton algorithm we discussed.
- No fault tolerance.
- Should be faster than our implementation:
 - More computational efficient: implemented in C++.
 - More communicational efficient: the slave-slave structure with all-reduce only communicates once per operation.



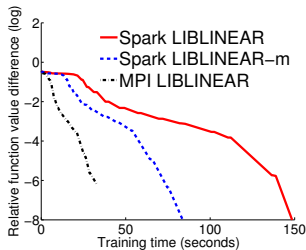
MPI LIBLINEAR

- A C++/MPI implementation by Zhuang et al. (2014) of the distributed trust region Newton algorithm we discussed.
- No fault tolerance.
- Should be faster than our implementation:
 - More computational efficient: implemented in C++.
 - More communicational efficient: the slave-slave structure with all-reduce only communicates once per operation.
- Should be faster, but need to know how large is the difference.

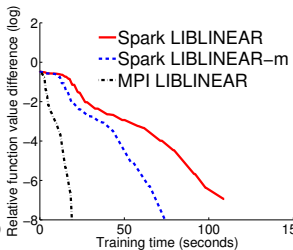


Spark versus MPI

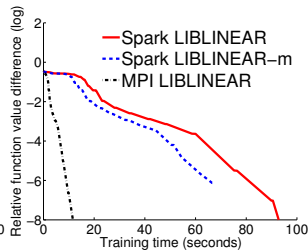
2 nodes



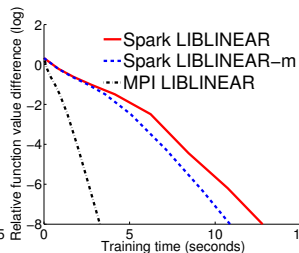
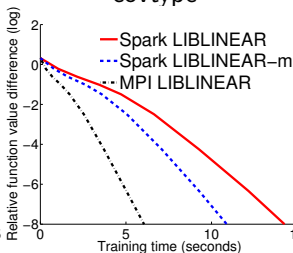
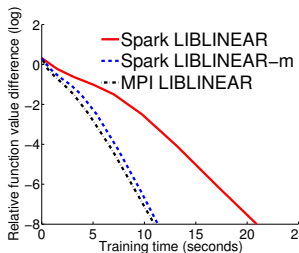
4 nodes



8 nodes



covtype

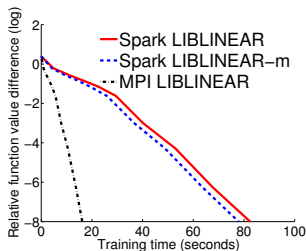


webspam

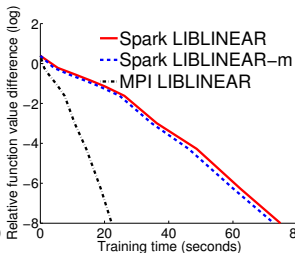


Spark versus MPI (Cont'd)

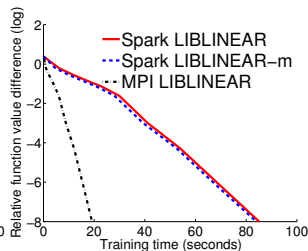
2 nodes



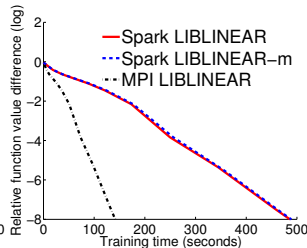
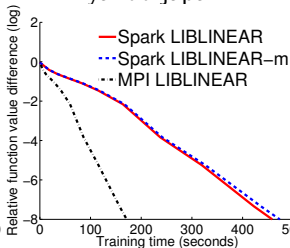
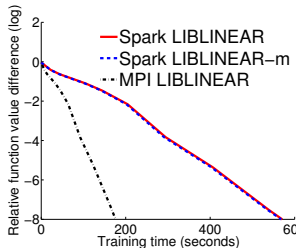
4 nodes



8 nodes



yahoo-japan

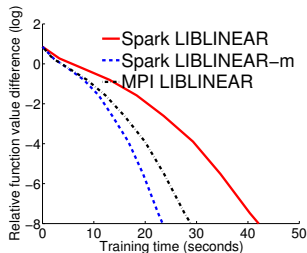


yahoo-korea

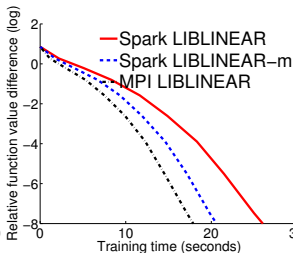


Spark versus MPI (Cont'd)

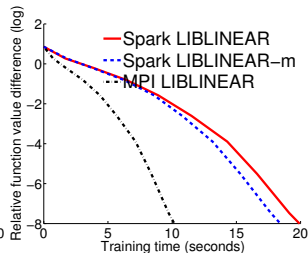
2 nodes



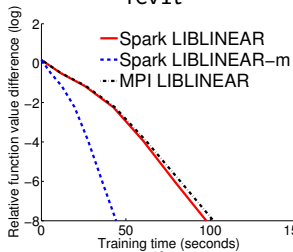
4 nodes



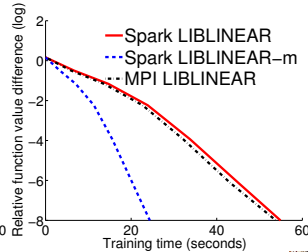
8 nodes



rcv1t



epsilon



Outline

- 1 Introduction
- 2 Our approach
- 3 Implementation design
- 4 Related Works
- 5 Discussions and Conclusions



Weakness and Future Work

- Integrating with MLlib (ongoing).



Weakness and Future Work

- Integrating with MLlib (ongoing).
- **Feature-wise approach** (Zhuang et al., 2014): communication cost can be reduced from $O(n)$ to $O(l)$ if $n \gg l$.



Weakness and Future Work

- Integrating with MLlib (ongoing).
- **Feature-wise approach** (Zhuang et al., 2014): communication cost can be reduced from $O(n)$ to $O(l)$ if $n \gg l$.
- Comparing with other newly available optimization approaches implemented on Spark (l-bfgs, dual coordinate ascent (Jaggi et al., 2014), etc.)



Conclusions

- We consider a distributed trust region Newton algorithm on Spark for training LR and linear SVM.
- Many implementation issues are thoroughly studied with careful empirical examinations.
- Our implementation on Spark is competitive with state-of-the-art packages.
- **Spark LIBLINEAR** is an distributed extension of **LIBLINEAR** and it is available at <http://www.csie.ntu.edu.tw/~cjlin/libsvmtools/distributed-liblinear/>.

